

# Introduction to SQL

## **1. What is SQL, and why is it essential in database management?**

**SQL (Structured Query Language)** is a standardized programming language used for managing and manipulating relational databases. It allows users to interact with a database to perform a wide variety of tasks.

### **Core Functions of SQL:**

1. **Data Querying** – Retrieve specific data using SELECT statements.
2. **Data Manipulation** – Add, update, or delete data using INSERT, UPDATE, and DELETE.
3. **Data Definition** – Define database schema using CREATE, ALTER, and DROP.
4. **Data Access Control** – Manage user access and permissions with GRANT and REVOKE.

---

### **Why SQL Is Essential in Database Management:**

#### **1. Standardized Language:**

- SQL is the industry standard adopted by most relational database systems (like MySQL, PostgreSQL, SQL Server, and Oracle), ensuring compatibility and consistency.

#### **2. Efficient Data Handling:**

- It allows users to quickly and efficiently retrieve and manipulate large volumes of data.

### **3. Relational Database Support:**

- SQL is specifically designed to work with relational databases, which organize data into tables with relationships.

### **4. Data Integrity and Accuracy:**

- SQL supports constraints and transactions, which ensure data consistency and integrity.

### **5. Accessibility for Multiple Roles:**

- Used by developers, data analysts, and database administrators (DBAs), making it a versatile tool across many IT roles.

### **6. Automation and Reporting:**

- SQL queries can be used in scripts and reports for business intelligence, automating repetitive data tasks.

---

In short, SQL is essential because it provides a powerful, efficient, and universally accepted way to interact with and manage structured data in databases.

## **2. Explain the difference between DBMS and RDBMS.**

The terms **DBMS (Database Management System)** and **RDBMS (Relational Database Management System)** are often used

interchangeably, but they refer to different types of systems with key distinctions.

---

## 1. DBMS (Database Management System)

A **DBMS** is software that allows users to **store, retrieve, and manage data** in databases. It is a general term that encompasses all types of database systems.

**Key Characteristics:**

- Can manage **structured, semi-structured, or unstructured** data.
  - Does **not enforce relationships** between tables (if tables exist at all).
  - Examples: **File systems, XML databases, NoSQL databases**, etc.
  - Focuses on individual data storage rather than relationships.
- 

## 2. RDBMS (Relational Database Management System)

An **RDBMS** is a specific type of DBMS that **stores data in a structured, tabular form** and enforces relationships among the data using keys (primary and foreign keys).

**Key Characteristics:**

- **Based on the relational model** (introduced by E.F. Codd).
- Data is stored in **tables (relations)**.
- Enforces **relationships** between tables.
- Supports **SQL** as the standard query language.

- Ensures **ACID properties** (Atomicity, Consistency, Isolation, Durability) for transaction reliability.
  - Examples: **MySQL, PostgreSQL, Oracle, Microsoft SQL Server.**
- 

### **Summary Comparison Table:**

<b>Feature</b>	<b>DBMS</b>	<b>RDBMS</b>
Data Storage	Any format	In structured tables
Data Relationships	Not enforced	Enforced via keys
Data Integrity	Basic or manual	Strong integrity with constraints
SQL Support	Optional or limited	Full SQL support
Examples	XML DB, File System, NoSQL	MySQL, Oracle, PostgreSQL
Transactions & ACID	Not always supported	Fully supported

---

### **In Short:**

- **Every RDBMS is a DBMS, but not every DBMS is an RDBMS.**
- If your application needs structured, relational data and robust integrity, **RDBMS is the better choice.**

### **3. Describe the role of SQL in managing relational databases.**

SQL (Structured Query Language) plays a **central role** in managing **relational databases**, serving as the primary language for all interactions with data stored in tables. Its design aligns perfectly with the **relational model**, allowing users to define, manipulate, control, and query structured data efficiently.

---

#### **Key Roles of SQL in Managing Relational Databases**

##### **1. Data Definition (DDL – Data Definition Language)**

SQL is used to define and modify the structure of database objects like tables, indexes, and schemas.

- Commands: CREATE, ALTER, DROP
- Example:
- ```
CREATE TABLE Employees (
    ID INT PRIMARY KEY,
    Name VARCHAR(50),
    Salary DECIMAL(10, 2)
);
```

##### **2. Data Manipulation (DML – Data Manipulation Language)**

SQL allows users to insert, update, delete, and retrieve data in relational tables.

- Commands: INSERT, UPDATE, DELETE, SELECT
- Example:

- SELECT Name, Salary FROM Employees WHERE Salary > 50000;

### **3. Data Querying**

SQL enables users to perform complex queries across multiple related tables using **joins**, **filters**, **aggregations**, etc.

- Example:
- SELECT e.Name, d.DepartmentName
- FROM Employees e
- JOIN Departments d ON e.DeptID = d.ID;

### **4. Data Control (DCL – Data Control Language)**

SQL helps control access to data using permissions and roles.

- Commands: GRANT, REVOKE
- Example:
- GRANT SELECT ON Employees TO Analyst;

### **5. Transaction Management (TCL – Transaction Control Language)**

SQL provides commands to manage transactions, ensuring data consistency and integrity.

- Commands: BEGIN, COMMIT, ROLLBACK
- Example:
- BEGIN;
- UPDATE Accounts SET Balance = Balance - 100 WHERE ID = 1;

- UPDATE Accounts SET Balance = Balance + 100 WHERE ID = 2;
  - COMMIT;
- 

## How SQL Supports Relational Database Principles:

- **Tabular Structure:** Works naturally with rows and columns.
  - **Primary and Foreign Keys:** Enforces relationships.
  - **Constraints:** Maintains data integrity (e.g., NOT NULL, UNIQUE, CHECK).
  - **Normalization:** Helps manage normalized data through efficient joins.
- 

## In Summary:

SQL is the **core interface** between users/applications and relational databases. It enables:

- Creating and managing data structures
- Inserting, querying, and updating data
- Enforcing relationships and integrity
- Controlling access and managing transactions

Without SQL, managing relational databases would be inefficient, inconsistent, and error-prone.

## 4. What are the key features of SQL?

SQL (Structured Query Language) has several **key features** that make it the standard and powerful tool for interacting with relational databases. Here's a breakdown of its most important features:

---

### Key Features of SQL

#### 1. Data Querying

- **Primary use:** Retrieve specific data from one or more tables.
  - **Command:** SELECT
  - **Example:**
  - SELECT Name, Age FROM Customers WHERE Age > 30;
- 

#### 2. Data Manipulation (DML)

- **Used to manage data inside tables.**
  - **Commands:** INSERT, UPDATE, DELETE
  - **Example:**
  - INSERT INTO Employees (Name, Salary) VALUES ('Alice', 60000);
- 

#### 3. Data Definition (DDL)

- **Used to define or modify database schema.**
- **Commands:** CREATE, ALTER, DROP
- **Example:**

- CREATE TABLE Products (
  - ID INT PRIMARY KEY,
  - Name VARCHAR(100),
  - Price DECIMAL(10, 2)
  - );
- 

#### 4. Data Control (DCL)

- **Manages permissions and access control.**
  - **Commands:** GRANT, REVOKE
  - **Example:**
  - GRANT SELECT ON Orders TO SalesUser;
- 

#### 5. Transaction Control (TCL)

- **Manages changes to data and ensures consistency.**
  - **Commands:** BEGIN, COMMIT, ROLLBACK
  - **Example:**
  - BEGIN;
  - UPDATE Accounts SET Balance = Balance - 100 WHERE ID = 1;
  - UPDATE Accounts SET Balance = Balance + 100 WHERE ID = 2;
  - COMMIT;
-

## 6. Relational Data Handling

- SQL supports relationships between tables using primary and foreign keys.
  - Allows for complex joins and subqueries.
- 

## 7. Built-in Functions

- SQL includes functions for:
    - Aggregation: SUM(), AVG(), COUNT()
    - String manipulation: CONCAT(), SUBSTRING()
    - Date/time operations: NOW(), DATEADD()
- 

## 8. Data Integrity and Constraints

- Enforces rules to maintain clean and accurate data.
  - Examples: NOT NULL, UNIQUE, CHECK, PRIMARY KEY, FOREIGN KEY
- 

## 9. Portability

- SQL is standardized and used across many database systems (e.g., MySQL, PostgreSQL, SQL Server, Oracle), with minor variations.
- 

## 10. Declarative Language

- SQL is **declarative**, meaning you specify *what* data you want, not *how* to get it.

- This simplifies code and makes queries easier to write and understand.
- 

## Summary Table

| Feature               | Description                               |
|-----------------------|-------------------------------------------|
| Querying              | Retrieve data using SELECT                |
| Data Manipulation     | Modify data with INSERT, UPDATE, DELETE   |
| Schema Definition     | Create/alter tables using CREATE, ALTER   |
| Access Control        | Manage permissions with GRANT, REVOKE     |
| Transaction Control   | Ensure consistency using COMMIT, ROLLBACK |
| Built-in Functions    | Use predefined functions for calculations |
| Relational Support    | Work with multiple related tables         |
| Integrity Constraints | Maintain valid and consistent data        |
| Portability           | Works across many DBMS platforms          |
| Declarative Syntax    | Focus on <i>what</i> not <i>how</i>       |

---

## In Summary:

SQL is a **versatile**, **powerful**, and **widely adopted** language for working with relational databases, offering robust tools for data management, integrity, security, and performance.

# SQL Syntax

## 1. What are the basic components of SQL syntax?

The **basic components of SQL syntax** are the building blocks that make up SQL statements. Understanding these components helps in writing clear and effective queries.

---

### 1. Clauses

Clauses define the structure of SQL statements.

- **SELECT** – Specifies the columns to retrieve
  - SELECT Name, Age
- **FROM** – Specifies the table(s) involved
  - FROM Customers
- **WHERE** – Filters rows based on conditions
  - WHERE Age > 30
- **ORDER BY** – Sorts results
  - ORDER BY Name ASC
- **GROUP BY** – Groups rows that share values
  - GROUP BY Department
- **HAVING** – Filters groups after aggregation
  - HAVING COUNT(\*) > 5

---

### 2. Statements

Statements perform specific actions.

- **Data Querying:**
  - `SELECT * FROM Products;`
  - **Data Insertion:**
  - `INSERT INTO Employees (Name, Salary) VALUES ('Alice', 60000);`
  - **Data Update:**
  - `UPDATE Employees SET Salary = 65000 WHERE Name = 'Alice';`
  - **Data Deletion:**
  - `DELETE FROM Employees WHERE Name = 'Alice';`
  - **Table Creation:**
  - `CREATE TABLE Customers (ID INT, Name VARCHAR(50));`
- 

### 3. Expressions

Expressions perform calculations or evaluate to values.

- Arithmetic:
  - `SELECT Price * Quantity AS TotalCost FROM Orders;`
  - Logical:
  - `WHERE Age > 25 AND City = 'New York'`
- 

### 4. Predicates

Predicates return a boolean result (TRUE/FALSE/UNKNOWN) for conditions.

- =, <>, <, >, <=, >=
  - BETWEEN, IN, LIKE, IS NULL
  - EXISTS, ANY, ALL
- 

## 5. Operators

Used to combine or compare expressions.

- **Arithmetic Operators:** +, -, \*, /
  - **Logical Operators:** AND, OR, NOT
  - **Comparison Operators:** =, <>, <, >
- 

## 6. Keywords

SQL uses **reserved keywords** that define operations and structure, like:

- SELECT, FROM, WHERE, JOIN, CREATE, DELETE, UPDATE, INSERT, etc.
- 

## 7. Identifiers

These represent database objects like:

- Table names: Employees
  - Column names: Salary, Name
  - Aliases: AS e in FROM Employees AS e
- 

## 8. Literals

Fixed data values in SQL.

- Strings: 'John'
  - Numbers: 42, 3.14
  - Dates: '2025-07-16'
- 

### **Example: Full SQL Statement**

```
SELECT Name, Salary * 1.1 AS NewSalary  
FROM Employees  
WHERE Department = 'Sales' AND Salary > 50000  
ORDER BY NewSalary DESC;
```

---

### **Summary:**

#### **Component Purpose**

**Clauses** Structure the query (e.g., SELECT, WHERE)

**Statements** Perform operations (e.g., INSERT, UPDATE)

**Expressions** Perform calculations or logic

**Predicates** Test conditions for filtering

**Operators** Combine expressions and conditions

**Keywords** Define SQL syntax and behavior

**Identifiers** Refer to database objects

**Literals** Represent constant values

These components work together to form valid, readable, and powerful SQL queries.

## 2. Write the general structure of an SQL SELECT statement.

The **general structure** of an SQL SELECT statement follows a specific order of clauses, each serving a unique purpose in querying data from a relational database.

---

### General SQL SELECT Statement Structure

SELECT column1, column2, ...

FROM table\_name

[WHERE condition]

[GROUP BY column]

[HAVING group\_condition]

[ORDER BY column [ASC|DESC]]

[LIMIT number]; -- (or TOP in some systems)

---

### Explanation of Clauses:

| Clause | Description |
|--------|-------------|
|--------|-------------|

**SELECT**      Specifies the columns to retrieve (use \* to select all)

**FROM**        Specifies the table(s) from which to retrieve data

| <b>Clause</b>      | <b>Description</b>                                            |
|--------------------|---------------------------------------------------------------|
| <b>WHERE</b>       | Filters rows based on a condition (optional)                  |
| <b>GROUP BY</b>    | Groups rows that share a common value (used with aggregation) |
| <b>HAVING</b>      | Filters groups after GROUP BY (like WHERE, but for groups)    |
| <b>ORDER BY</b>    | Sorts the result set (ascending ASC or descending DESC)       |
| <b>LIMIT / TOP</b> | Limits the number of rows returned (optional)                 |

---

### **Example:**

```

SELECT Department, COUNT(*) AS EmployeeCount
FROM Employees
WHERE Status = 'Active'
GROUP BY Department
HAVING COUNT(*) > 5
ORDER BY EmployeeCount DESC
LIMIT 10;

```

---

### **Notes:**

- The order of clauses **must** be maintained as shown above.
- Some databases (like SQL Server) use TOP instead of LIMIT:

- `SELECT TOP 10 * FROM Employees;`
  - Aliases (AS) can be used to rename columns or tables for readability.
- 

This structure forms the foundation for most SQL data queries.

### **3. Explain the role of clauses in SQL statements.**

In SQL, **clauses** are essential components that define the **structure and behavior** of a statement. Each clause serves a specific purpose in helping you retrieve, filter, group, sort, or limit data from a database.

---

#### **What Are Clauses in SQL?**

Clauses are **keywords** or **parts of SQL statements** that perform specific functions in query construction. A full SQL statement is often made up of multiple clauses combined in a logical sequence.

---

#### **Common SQL Clauses and Their Roles:**

| <b>Clause</b> | <b>Purpose</b>                                                       |
|---------------|----------------------------------------------------------------------|
| <b>SELECT</b> | Specifies the columns (or expressions) to be retrieved from a table. |
| <b>FROM</b>   | Identifies the table(s) from which to retrieve data.                 |

| <b>Clause</b>      | <b>Purpose</b>                                                                |
|--------------------|-------------------------------------------------------------------------------|
| <b>WHERE</b>       | Filters rows based on a condition before grouping or aggregation.             |
| <b>GROUP BY</b>    | Groups rows that have the same values in specified columns.                   |
| <b>HAVING</b>      | Filters grouped results (used with GROUP BY and aggregate functions).         |
| <b>ORDER BY</b>    | Sorts the result set by one or more columns in ascending or descending order. |
| <b>LIMIT / TOP</b> | Limits the number of rows returned by the query (syntax varies by database).  |

---

### **Example SQL Statement Using Multiple Clauses:**

```

SELECT Department, COUNT(*) AS EmployeeCount
FROM Employees
WHERE Status = 'Active'
GROUP BY Department
HAVING COUNT(*) > 10
ORDER BY EmployeeCount DESC
LIMIT 5;

```

### **Explanation:**

- **SELECT:** Chooses columns to display.
- **FROM:** Indicates the source table.

- **WHERE:** Filters to include only active employees.
  - **GROUP BY:** Groups employees by department.
  - **HAVING:** Includes only departments with more than 10 active employees.
  - **ORDER BY:** Sorts departments by employee count (highest first).
  - **LIMIT:** Returns only the top 5 departments.
- 

### Why Clauses Matter:

- They provide **modularity**: each clause handles one aspect of the query.
  - They allow **fine-grained control** over data retrieval.
  - Using the correct **order** of clauses is **crucial** for valid SQL syntax.
- 

### In Summary:

Clauses are the **backbone of SQL statements**, each playing a distinct role in shaping the query logic—from selecting data and filtering rows to grouping and sorting results. Mastering clauses is essential to writing effective and accurate SQL queries.

# SQL Constraints

## 1. What are constraints in SQL? List and explain the different types of constraints.

In **SQL**, **constraints** are rules applied to columns or tables to enforce **data integrity**, **accuracy**, and **consistency** within a relational database.

They ensure that only valid data is entered and maintained in the database, preventing errors and preserving the quality of stored information.

---

### Types of Constraints in SQL (with Explanation):

#### Constraint Description

##### 1. NOT NULL

- Ensures that a column **cannot have NULL values**.
  - Used when a value is **required** for every row.
  - Example:
  - `CREATE TABLE Employees (`
  - `ID INT NOT NULL,`
  - `Name VARCHAR(50) NOT NULL`
  - `);`
- 

##### 2. UNIQUE

- Ensures that **all values in a column are different**.

- Allows **NULL** (unless combined with NOT NULL).
  - Example:
  - CREATE TABLE Users (
  - Email VARCHAR(100) UNIQUE
  - );
- 

### 3. PRIMARY KEY

- Uniquely identifies each row in a table.
  - Combines the behavior of NOT NULL and UNIQUE.
  - A table can have **only one** primary key (can be a composite key).
  - Example:
  - CREATE TABLE Students (
  - StudentID INT PRIMARY KEY,
  - Name VARCHAR(100)
  - );
- 

### 4. FOREIGN KEY

- Establishes a **relationship between two tables**.
- Ensures that the value in one table matches a value in another (parent) table.
- Enforces **referential integrity**.
- Example:

- CREATE TABLE Orders (
  - OrderID INT PRIMARY KEY,
  - CustomerID INT,
  - FOREIGN KEY (CustomerID) REFERENCES  
Customers(CustomerID)
  - );
- 

## 5. CHECK

- Ensures that values in a column **meet a specific condition.**
  - Adds a custom validation rule.
  - Example:
  - CREATE TABLE Products (
  - Price DECIMAL(10, 2),
  - CHECK (Price > 0)
  - );
- 

## 6. DEFAULT

- Provides a **default value** for a column when no value is specified.
- Helps reduce NULLs and ensures consistent input.
- Example:
- CREATE TABLE Accounts (
- Balance DECIMAL(10, 2) DEFAULT 0.00

- );
- 

### **Summary Table:**

| <b>Constraint</b> | <b>Purpose</b>                                        |
|-------------------|-------------------------------------------------------|
| NOT NULL          | Prevents NULL values in a column                      |
| UNIQUE            | Ensures all values in a column are unique             |
| PRIMARY KEY       | Uniquely identifies a row (implies NOT NULL + UNIQUE) |
| FOREIGN KEY       | Maintains referential integrity between tables        |
| CHECK             | Validates data based on a condition                   |
| DEFAULT           | Sets a default value for a column                     |

---

### **In Practice:**

Using constraints helps enforce business rules **at the database level**, reducing the risk of incorrect or inconsistent data even before it reaches the application layer.

## **2. How do PRIMARY KEY and FOREIGN KEY constraints differ?**

The **PRIMARY KEY** and **FOREIGN KEY** constraints are both essential in SQL for maintaining **data integrity**, but they serve **different purposes** within a relational database.

---

## PRIMARY KEY vs FOREIGN KEY – Key Differences

| Feature                 | PRIMARY KEY                                              | FOREIGN KEY                                                         |
|-------------------------|----------------------------------------------------------|---------------------------------------------------------------------|
| <b>Definition</b>       | Uniquely identifies each row in a table                  | Links one table to another, enforcing referential integrity         |
| <b>Uniqueness</b>       | Must be <b>unique</b> for each row                       | Can contain <b>duplicate</b> values                                 |
| <b>NULL values</b>      | Cannot be <b>NULL</b> (implies NOT NULL)                 | Can be <b>NULL</b> unless explicitly restricted                     |
| <b>Location</b>         | Defined <b>in the current table</b>                      | Points to the <b>primary key (or unique key) of another table</b>   |
| <b>Purpose</b>          | Acts as the <b>main identifier</b> for a record          | Ensures that a value exists in the <b>referenced (parent) table</b> |
| <b>How many allowed</b> | Only <b>one primary key</b> per table (can be composite) | Can have <b>multiple foreign keys</b> per table                     |

---

### Example:

#### Customers table (with PRIMARY KEY):

```
CREATE TABLE Customers (
```

```
    CustomerID INT PRIMARY KEY,
```

```
    Name VARCHAR(100)
```

);

### **Orders table (with FOREIGN KEY):**

```
CREATE TABLE Orders (
    OrderID INT PRIMARY KEY,
    CustomerID INT,
    OrderDate DATE,
    FOREIGN KEY (CustomerID) REFERENCES
    Customers(CustomerID)
);
```

- Customers.CustomerID is the **PRIMARY KEY** – each customer must have a unique ID.
- Orders.CustomerID is a **FOREIGN KEY** – it must match a valid CustomerID from the Customers table or be NULL.

---

### **In Summary:**

- A **PRIMARY KEY** enforces **uniqueness and identity** in its own table.
- A **FOREIGN KEY** enforces **relationships and consistency** between two tables.

Together, they help establish the **relational structure** that defines a relational database.

### **3. What is the role of NOT NULL and UNIQUE constraints?**

The **NOT NULL** and **UNIQUE** constraints in SQL are used to enforce **data integrity** at the column level by controlling the type of data that can be stored.

---

#### **1. NOT NULL Constraint**

##### **Role:**

- Ensures that a column **cannot contain NULL values**.
- Used when a value is **required** in every row for a particular column.

##### **Use Case:**

- Prevents missing or unknown values in critical fields like usernames, email addresses, or IDs.

##### **Example:**

```
CREATE TABLE Employees (
    EmployeeID INT NOT NULL,
    Name VARCHAR(100) NOT NULL
);
```

- Every EmployeeID and Name **must** be provided — cannot be left blank.
- 

#### **2. UNIQUE Constraint**

##### **Role:**

- Ensures that **all values in a column (or combination of columns) are distinct.**
- Allows only **one occurrence of each value** in that column.

### Use Case:

- Enforces **data uniqueness** for fields like email, usernames, or license numbers.

### Example:

```
CREATE TABLE Users (
    Email VARCHAR(100) UNIQUE
);
```

- No two users can register with the **same email address**.

### Comparison Table

| Constraint Enforces | Allows NULL?      | Purpose                                |
|---------------------|-------------------|----------------------------------------|
| NOT NULL            | Value is required | Ensures that a field must have a value |
| UNIQUE              | Value is unique   | Prevents duplicate values in a column  |

\* UNIQUE allows multiple NULLs (depending on the SQL dialect), because NULL is treated as "unknown", not equal to anything — including other NULLs.

### In Practice:

- You can combine both constraints on the same column:
- Email VARCHAR(100) UNIQUE NOT NULL

This means **each email must be present and unique** — ideal for login systems or user IDs.

---

### **Summary:**

- **NOT NULL**: Makes sure data **is provided**.
- **UNIQUE**: Makes sure data **is not duplicated**.

Together, they are powerful tools for maintaining **clean, reliable, and meaningful data**.

# Main SQL Commands and Sub-commands (DDL)

## 1. Define the SQL Data Definition Language (DDL).

### **SQL Data Definition Language (DDL)**

**Data Definition Language (DDL)** is a category of SQL statements used to **define and manage the structure of database objects**, such as tables, schemas, indexes, and views.

---

#### **Key Functions of DDL:**

DDL allows you to:

- **Create** new database objects
- **Alter** existing structures
- **Drop** (delete) objects
- **Rename** or **truncate** data within structures

DDL deals with the **schema** (structure), not the data itself.

---

#### **Common DDL Commands:**

##### **Command Description**

**CREATE**      Creates a new database, table, view, index, etc.

**ALTER**      Modifies an existing database object (e.g., add a column)

**DROP**      Deletes a database object permanently

## **Command Description**

|                 |                                                                                  |
|-----------------|----------------------------------------------------------------------------------|
| <b>TRUNCATE</b> | Removes all records from a table <b>without logging</b> individual row deletions |
| <b>RENAME</b>   | Renames a table or column (supported in some DBMS)                               |

---

## **Examples of DDL Commands:**

### **CREATE Table**

```
CREATE TABLE Employees (
    EmployeeID INT PRIMARY KEY,
    Name VARCHAR(100) NOT NULL,
    Salary DECIMAL(10, 2)
);
```

### **ALTER Table (Add Column)**

```
ALTER TABLE Employees
ADD HireDate DATE;
```

### **DROP Table**

```
DROP TABLE Employees;
```

### **TRUNCATE Table**

```
TRUNCATE TABLE Employees;
```

---

## **DDL and Transactions:**

- Most DDL operations are **auto-committed** — changes are **immediate and irreversible** without a rollback.
- 

### In Summary:

| Feature              | DDL Description                                                                                                  |
|----------------------|------------------------------------------------------------------------------------------------------------------|
| Purpose              | Define and modify database structure                                                                             |
| Affects              | Tables, schemas, views, indexes, etc.                                                                            |
| Common Commands      | CREATE, ALTER, DROP, TRUNCATE                                                                                    |
| Transaction Behavior | Usually auto-committed                                                                                           |
| <b>DDL</b>           | is foundational to setting up a relational database — it provides the blueprint for storing and organizing data. |

## 2. Explain the **CREATE** command and its syntax.

### SQL CREATE Command

The **CREATE** command in SQL is a **Data Definition Language (DDL)** statement used to **create new database objects** such as:

- Tables
  - Databases
  - Views
  - Indexes
  - Schemas
  - Stored procedures (in some DBMSs)
-

## **1. Purpose of CREATE**

- Define the **structure** (schema) of a database object.
  - Specify **columns, data types, constraints, and relationships**.
- 

## **2. Basic Syntax: Creating a Table**

```
CREATE TABLE table_name (
    column1 datatype [constraint],
    column2 datatype [constraint],
    ...
);
```

---

## **3. Example: Creating a Table with Constraints**

```
CREATE TABLE Employees (
    EmployeeID INT PRIMARY KEY,
    FirstName VARCHAR(50) NOT NULL,
    LastName VARCHAR(50),
    Email VARCHAR(100) UNIQUE,
    HireDate DATE,
    Salary DECIMAL(10, 2) CHECK (Salary > 0),
    DepartmentID INT
);
```

### **Key Elements:**

- EmployeeID is the **primary key**
  - FirstName cannot be NULL
  - Email must be **unique**
  - Salary must be **greater than 0**
  - DepartmentID could later be set as a **foreign key**
- 

## 4. Other CREATE Syntax Examples

### **Create a Database:**

```
CREATE DATABASE CompanyDB;
```

### **Create a View:**

```
CREATE VIEW HighEarners AS  
SELECT FirstName, LastName, Salary  
FROM Employees  
WHERE Salary > 100000;
```

### **Create an Index:**

```
CREATE INDEX idx_lastname  
ON Employees (LastName);
```

---

## **Important Notes:**

- Most CREATE operations are **auto-committed** (cannot be rolled back).
- Column **constraints** (e.g., NOT NULL, UNIQUE, CHECK) improve data integrity.

- Use IF NOT EXISTS (if supported by your DBMS) to avoid errors:
  - CREATE TABLE IF NOT EXISTS Employees (...);
- 

### **Summary Table:**

| <b>Element</b>  | <b>Description</b>                              |
|-----------------|-------------------------------------------------|
| CREATE TABLE    | Defines a new table and its columns             |
| CREATE DATABASE | Creates a new database                          |
| CREATE VIEW     | Creates a virtual table based on a SELECT query |
| CREATE INDEX    | Speeds up queries on one or more columns        |

The CREATE command is your **first step in building database structure**, allowing you to define how and what kind of data will be stored.

### **3. What is the purpose of specifying data types and constraints during table creation?**

Specifying **data types** and **constraints** during table creation in SQL is essential for ensuring that the database is **structured correctly, efficient, and data-consistent** from the beginning.

---

#### **Purpose of Data Types**

**Data types** define the **kind of values** a column can store.

## Why Specify Data Types?

| Benefit                         | Explanation                                                                 |
|---------------------------------|-----------------------------------------------------------------------------|
| <b>Data integrity</b>           | Ensures only valid data is entered (e.g., INT for age, DATE for birthdate). |
| <b>Performance optimization</b> | Helps the database engine allocate memory and storage efficiently.          |
| <b>Enforces business logic</b>  | Prevents input errors (e.g., storing text in numeric fields).               |
| <b>Reduces data anomalies</b>   | Prevents storing inconsistent or ambiguous values.                          |

### Example:

Salary DECIMAL(10, 2) -- only allows numeric values with 2 decimal places

---

## Purpose of Constraints

**Constraints** enforce **rules on data** to maintain **accuracy** and **reliability** within the database.

## Why Use Constraints?

| Constraint  | Purpose                                                   |
|-------------|-----------------------------------------------------------|
| NOT NULL    | Ensures a column always has a value (no blanks).          |
| UNIQUE      | Prevents duplicate values in a column.                    |
| PRIMARY KEY | Uniquely identifies each row (implies NOT NULL + UNIQUE). |

## **Constraint      Purpose**

|                |                                               |
|----------------|-----------------------------------------------|
| FOREIGN<br>KEY | Maintains valid relationships between tables. |
| CHECK          | Enforces custom rules (e.g., salary > 0).     |
| DEFAULT        | Sets a default value if none is provided.     |

---

## **Example: Table with Data Types and Constraints**

```
CREATE TABLE Employees (
    EmployeeID INT PRIMARY KEY,
    Name VARCHAR(100) NOT NULL,
    Email VARCHAR(100) UNIQUE,
    Salary DECIMAL(10, 2) CHECK (Salary > 0),
    JoinDate DATE DEFAULT CURRENT_DATE
);
```

---

## **In Summary:**

### **Feature      Purpose**

|             |                                                                               |
|-------------|-------------------------------------------------------------------------------|
| Data Types  | Define <b>what kind of data</b> each column stores (e.g., INT, VARCHAR, DATE) |
| Constraints | Enforce <b>rules</b> to keep data accurate, unique, and meaningful            |

Together, they help create a **robust, clean, and reliable** database structure right from the start.

# ALTER Command

## 1. What is the use of the ALTER command in SQL?

The ALTER command in SQL is used to **modify the structure of an existing database object**, such as a table. It allows you to change a table's schema without having to delete and recreate the table.

### Common Uses of ALTER in SQL:

#### 1. Add a New Column:

2. ALTER TABLE employees ADD birthdate DATE;

#### 3. Drop a Column:

4. ALTER TABLE employees DROP COLUMN birthdate;

#### 5. Modify a Column (e.g., data type or size):

6. ALTER TABLE employees MODIFY salary DECIMAL(10, 2);

(Note: Some databases use *MODIFY*, others use *ALTER COLUMN*.)

#### 7. Rename a Column (syntax varies by database):

8. -- In PostgreSQL

9. ALTER TABLE employees RENAME COLUMN fullname TO name;

10.

11. -- In MySQL

12. ALTER TABLE employees CHANGE fullname name VARCHAR(100);

#### 13. Rename a Table:

14. ALTER TABLE employees RENAME TO staff;
15. **Add/Drop Constraints** (e.g., primary key, foreign key, unique):
16. ALTER TABLE employees ADD CONSTRAINT pk\_emp PRIMARY KEY (id);
17. ALTER TABLE employees DROP CONSTRAINT pk\_emp;

### **Summary:**

The ALTER command is essential for **maintaining and evolving** the structure of a database as requirements change over time.

## **2. How can you add, modify, and drop columns from a table using ALTER?**

You can use the ALTER TABLE statement in SQL to **add, modify, and drop** columns in an existing table. Here's how you do each operation:

---

### **1. Add a Column**

Use the ADD keyword:

```
ALTER TABLE table_name
```

```
ADD column_name data_type;
```

#### **Example:**

```
ALTER TABLE employees
```

```
ADD birthdate DATE;
```

You can also add multiple columns at once:

```
ALTER TABLE employees  
ADD (middle_name VARCHAR(50), hire_date DATE);
```

---

## 2. Modify a Column

Use the MODIFY or ALTER COLUMN keyword (depends on the database system):

- **MySQL / Oracle:**
  - ALTER TABLE employees
  - MODIFY birthdate DATETIME;
- **SQL Server:**
  - ALTER TABLE employees
  - ALTER COLUMN birthdate DATETIME;
- **PostgreSQL:**
  - ALTER TABLE employees
  - ALTER COLUMN birthdate TYPE TIMESTAMP;

---

## 3. Drop a Column

Use the DROP COLUMN keyword:

```
ALTER TABLE table_name  
DROP COLUMN column_name;
```

### **Example:**

```
ALTER TABLE employees  
DROP COLUMN middle_name;
```

You can drop multiple columns in some databases like Oracle or PostgreSQL:

```
ALTER TABLE employees  
DROP COLUMN middle_name,  
DROP COLUMN hire_date;
```

---

### **Notes:**

- Always **back up your data** before altering tables, especially when dropping columns.
- Syntax may vary slightly depending on the SQL dialect (MySQL, PostgreSQL, SQL Server, Oracle, etc.).

# DROP Command

## 1. What is the function of the DROP command in SQL?

The DROP command in SQL is used to **permanently delete database objects**, such as **tables**, **databases**, **views**, **indexes**, or **constraints**. Once an object is dropped, **all data and structure are lost** and **cannot be recovered** unless you have a backup.

---

### **Common Uses of the DROP Command:**

#### **1. Drop a Table**

Removes the table and all its data.

`DROP TABLE employees;`

#### **2. Drop a Database**

Deletes the entire database and all its objects (tables, views, etc.).

`DROP DATABASE company_db;`

#### **3. Drop a Column Constraint**

Removes a constraint like PRIMARY KEY, FOREIGN KEY, or UNIQUE.

- Drop a constraint (example in PostgreSQL or SQL Server):
- `ALTER TABLE employees DROP CONSTRAINT emp_pk;`

#### **4. Drop an Index**

Deletes an index:

```
DROP INDEX idx_name ON employees;
```

(*Syntax varies by DBMS; in SQL Server: DROP INDEX idx\_name  
ON employees; in MySQL: DROP INDEX idx\_name ON  
employees;*)

## 5. Drop a View

```
DROP VIEW employee_summary;
```

---

### Important Notes:

- The **DROP** command **cannot be undone** — use it with caution.
  - Any object dropped using **DROP** will no longer be accessible.
  - Dropping a table also drops all related indexes, triggers, and constraints.
- 

### Summary:

| Command         | Effect                                    |
|-----------------|-------------------------------------------|
| DROP TABLE      | Deletes a table and its data              |
| DROP DATABASE   | Deletes the database entirely             |
| DROP CONSTRAINT | Removes a key or rule from a table column |
| DROP INDEX      | Deletes an index used for fast lookup     |
| DROP VIEW       | Deletes a stored SQL query view           |

## 2. What are the implications of dropping a table from a database?

Dropping a table from a database using the DROP TABLE command has **serious and permanent implications**. Here's what happens and what to consider before using it:

---

### Implications of Dropping a Table

#### 1. Permanent Data Loss

- All rows (records) in the table are permanently deleted.
- This action **cannot be undone** (unless you have a backup).

#### 2. Loss of Structure

- The **table definition** (columns, data types, constraints) is also removed.
- You'll have to recreate it manually if needed later.

#### 3. Cascading Effects (in some cases)

- If the table has **foreign key relationships**, other tables may be affected:
  - The DROP may fail unless those dependencies are removed first.
  - Or, with CASCADE (in PostgreSQL, for example), related constraints and objects will also be dropped.
- DROP TABLE orders CASCADE;

#### 4. Loss of Related Objects

- Indexes, triggers, and constraints **associated with the table** are also deleted automatically.

## 5. Application Errors

- Any applications or queries that rely on the dropped table will fail.
  - This can lead to runtime errors and system instability.
- 

## Best Practices Before Dropping a Table

### 1. Back Up Data:

- Always create a backup (export to a file, dump, etc.) before dropping.

### 2. Check Dependencies:

- Use system views or metadata tools to identify foreign keys, views, or procedures that use the table.

### 3. Inform Stakeholders:

- Make sure your team or application owners know you're dropping the table.

### 4. Use a Staging Environment:

- Test the impact in a development or staging environment before applying in production.
- 

## Summary

| Impact    | Description                         |
|-----------|-------------------------------------|
| Data Loss | All records are permanently deleted |

| Impact                  | Description                                    |
|-------------------------|------------------------------------------------|
| Schema Loss             | Table structure and constraints are removed    |
| Dependency Risks        | May break other tables, views, or applications |
| No Undo                 | Irreversible unless restored from a backup     |
| Associated Objects Lost | Indexes, triggers, constraints are deleted too |

---

If you're unsure whether to drop a table, consider alternatives like:

- TRUNCATE TABLE (clears data but keeps structure)
- RENAME TABLE (temporarily hide or archive)
- CREATE TABLE new\_table AS SELECT \* FROM old\_table  
(backup clone)

# Data Manipulation Language (DML)

## 1. Define the INSERT, UPDATE, and DELETE commands in SQL.

In SQL, the INSERT, UPDATE, and DELETE commands are part of **Data Manipulation Language (DML)**. These commands are used to manage the data stored within tables.

---

### 1. INSERT – Add New Records

The INSERT command is used to **add new rows** to a table.

#### Syntax:

```
INSERT INTO table_name (column1, column2, ...)  
VALUES (value1, value2, ...);
```

#### Example:

```
INSERT INTO employees (id, name, department)  
VALUES (101, 'Alice', 'HR');
```

---

### 2. UPDATE – Modify Existing Records

The UPDATE command is used to **change existing data** in one or more rows of a table.

#### Syntax:

```
UPDATE table_name  
SET column1 = value1, column2 = value2, ...  
WHERE condition;
```

### **Example:**

```
UPDATE employees  
SET department = 'Finance'  
WHERE id = 101;
```

⚠ Always include a WHERE clause to avoid updating **all rows** unintentionally.

---

### **3. ✗ DELETE – Remove Records**

The DELETE command is used to **remove one or more rows** from a table.

#### **Syntax:**

```
DELETE FROM table_name  
WHERE condition;
```

#### **Example:**

```
DELETE FROM employees  
WHERE id = 101;
```

⚠ Omitting the WHERE clause will delete **all rows** in the table.

---

#### **✓ Summary Table:**

| <b>Command Purpose</b> | <b>Affects</b>         |
|------------------------|------------------------|
| INSERT                 | Adds new data          |
| UPDATE                 | Modifies existing data |
|                        | Changes existing rows  |

| Command Purpose          | Affects      |
|--------------------------|--------------|
| DELETE      Removes data | Deletes rows |

## 2. What is the importance of the WHERE clause in UPDATE and DELETE operations?

The WHERE clause is **critically important** in UPDATE and DELETE operations in SQL because it specifies **which rows** should be affected by the operation.

---

### Without WHERE: Dangerous!

- If you **omit the WHERE clause**, **all rows** in the table will be updated or deleted.

#### Example:

`DELETE FROM employees;`

This would **delete every record** in the employees table — probably not what you intended.

---

### With WHERE: Targeted and Safe

The WHERE clause helps you **narrow down** the affected rows based on specific conditions.

#### Example – Update:

`UPDATE employees`

`SET department = 'Finance'`

WHERE id = 101;

- ✓ Only updates the employee with id = 101.

### Example – Delete:

DELETE FROM employees

WHERE department = 'HR';

- ✓ Only deletes employees in the HR department.
- 

### Why It Matters:

| Reason                | Explanation                                            |
|-----------------------|--------------------------------------------------------|
| <b>Data Integrity</b> | Prevents accidental changes or loss of important data. |
| <b>Precision</b>      | Ensures only intended rows are affected.               |
| <b>Performance</b>    | Speeds up queries by avoiding full table scans.        |
| <b>Safety</b>         | Reduces the risk of catastrophic errors.               |

---

### Pro Tip:

Always **double-check your WHERE clause** before running an UPDATE or DELETE. If unsure, try running a SELECT with the same condition first to preview the affected rows:

SELECT \* FROM employees WHERE department = 'HR';

# Data Query Language (DQL)

## 1. What is the SELECT statement, and how is it used to query data?

The SELECT statement is the **most commonly used SQL command**. It is used to **query and retrieve data** from one or more tables in a database.

---

### What is the SELECT Statement?

SELECT allows you to **specify exactly what data** you want to view — including which columns, from which tables, and under what conditions.

---

### Basic Syntax

SELECT column1, column2, ...

FROM table\_name

WHERE condition;

### Example:

SELECT name, department

FROM employees

WHERE department = 'HR';

✓ This returns the names and departments of employees who work in HR.

---

## Common Clauses Used with SELECT

| Clause      | Description                                  |
|-------------|----------------------------------------------|
| FROM        | Specifies the table(s) to retrieve data from |
| WHERE       | Filters rows based on a condition            |
| ORDER BY    | Sorts the result set                         |
| GROUP BY    | Groups rows that have the same values        |
| HAVING      | Filters groups created by GROUP BY           |
| JOIN        | Combines rows from multiple tables           |
| LIMIT / TOP | Restricts number of returned rows            |

---

## Examples

### 1. Select All Columns

```
SELECT * FROM employees;
```

### 2. Select Specific Columns

```
SELECT name, salary FROM employees;
```

### 3. With WHERE Clause

```
SELECT name FROM employees WHERE salary > 50000;
```

### 4. With ORDER BY

```
SELECT name, salary FROM employees ORDER BY salary DESC;
```

### 5. With LIMIT (MySQL/PostgreSQL)

```
SELECT * FROM employees LIMIT 5;
```

---

## **Summary**

### **Feature      Purpose**

**SELECT**      Choose what data to retrieve

**FROM**      Specify the table(s)

**WHERE**      Filter rows based on conditions

**ORDER BY** Sort the result

**LIMIT**      Limit number of results returned

## **2. Explain the use of the ORDER BY and WHERE clauses in SQL queries.**

The WHERE and ORDER BY clauses in SQL are used to **filter** and **sort** data in SELECT (and other DML) queries. They help you get **precise and organized results** from your database.

---

### **1. WHERE Clause – Filters Rows**

#### **Purpose:**

The WHERE clause is used to **filter records** based on specified conditions. Only rows that meet the condition are included in the result set.

#### **Syntax:**

```
SELECT column1, column2
```

```
FROM table_name
```

```
WHERE condition;
```

### **Example:**

```
SELECT name, department  
FROM employees  
WHERE department = 'HR';
```

Returns only employees in the HR department.

### **Without WHERE:**

```
SELECT * FROM employees;
```

Returns **all** employees, regardless of department.

---

## **2. ORDER BY Clause – Sorts Rows**

### **Purpose:**

The ORDER BY clause is used to **sort the result set** by one or more columns in **ascending (ASC)** or **descending (DESC)** order.

### **Syntax:**

```
SELECT column1, column2  
FROM table_name  
ORDER BY column1 [ASC|DESC];
```

### **Example:**

```
SELECT name, salary  
FROM employees  
ORDER BY salary DESC;
```

Returns employees sorted by salary, highest to lowest.

---

## Using WHERE and ORDER BY Together

```
SELECT name, salary  
FROM employees  
WHERE department = 'Sales'  
ORDER BY salary DESC;
```

Filters for only Sales employees and sorts them by salary in descending order.

---

## Summary Table

| Clause   | Function                        | Notes                           |
|----------|---------------------------------|---------------------------------|
| WHERE    | Filters rows based on condition | Comes <b>before</b> ORDER BY    |
| ORDER BY | Sorts rows in the result set    | Default is ASC if not specified |

---

# Data Control Language (DCL)

## 1. What is the purpose of GRANT and REVOKE in SQL?

The GRANT and REVOKE commands in SQL are used to **manage user permissions and access control** in a database. They allow database administrators to control **who can do what** with database objects (like tables, views, procedures, etc.).

---

### 1. GRANT – Give Permissions

#### Purpose:

The GRANT command is used to **give users specific privileges** to perform actions on database objects.

#### Syntax:

GRANT privilege [, ...]

ON object\_name

TO user\_name;

#### Example:

GRANT SELECT, INSERT

ON employees

TO user1;

Grants user1 the ability to **view and add data** to the employees table.

---

## **2. REVOKE – Remove Permissions**

### **Purpose:**

The REVOKE command is used to **take back privileges** that were previously granted to users.

### **Syntax:**

```
REVOKE privilege [, ...]
```

```
    ON object_name
```

```
    FROM user_name;
```

### **Example:**

```
REVOKE INSERT
```

```
    ON employees
```

```
    FROM user1;
```

Removes the ability of user1 to insert data into the employees table.

---

### **Common Privileges in SQL:**

#### **Privilege Meaning**

**SELECT** Read data from a table or view

**INSERT** Add new rows

**UPDATE** Modify existing rows

**DELETE** Remove rows

**ALL** All available privileges

## Privilege Meaning

EXECUTE Run a stored procedure or function

---

### Summary:

| Command Purpose | Effect             |
|-----------------|--------------------|
| GRANT           | Give permissions   |
| REVOKE          | Remove permissions |

---

These commands are essential for **security**, **data integrity**, and **controlled access** in a multi-user database environment.

## 2. How do you manage privileges using these commands?

Managing privileges using the GRANT and REVOKE commands in SQL involves **assigning** or **removing** user access to perform specific actions on database objects like tables, views, or procedures.

Here's how you can **effectively manage privileges** step by step:

---

### 1. Using GRANT to Give Privileges

#### Syntax:

GRANT privilege [, ...]

ON object\_name

TO user\_name [, ...];

**Example – Granting Table Access:**

GRANT SELECT, INSERT, UPDATE

ON employees

TO user1;

user1 can now view, add, and modify data in the employees table.

**Example – Grant All Privileges:**

GRANT ALL

ON employees

TO user2;

**Example – Granting EXECUTE on Stored Procedure:**

GRANT EXECUTE

ON PROCEDURE update\_salary

TO user3;

---

## 2. Using REVOKE to Remove Privileges

**Syntax:**

REVOKE privilege [, ...]

ON object\_name

FROM user\_name [, ...];

**Example – Revoke Insert Privilege:**

REVOKE INSERT

```
ON employees
```

```
FROM user1;
```

user1 can no longer insert new rows into the employees table.

### **Example – Revoke All Privileges:**

```
REVOKE ALL
```

```
ON employees
```

```
FROM user2;
```

---

## **3. Managing Privileges for Roles (Optional Advanced Use)**

Instead of assigning privileges directly to individual users, you can assign them to **roles**, then grant roles to users.

### **Step 1: Create Role**

```
CREATE ROLE hr_role;
```

### **Step 2: Grant Privileges to Role**

```
GRANT SELECT, INSERT ON employees TO hr_role;
```

### **Step 3: Assign Role to User**

```
GRANT hr_role TO user1;
```

### **Benefit:**

Easier privilege management for multiple users performing similar tasks.

---

## **Best Practices**

- **Always use least privilege** — only grant what's necessary.

- **Use roles** for group-based permission management.
  - **Regularly audit** who has what privileges.
  - **Document changes** for accountability.
- 

## Summary

| Task                        | Command                      | Example                              |
|-----------------------------|------------------------------|--------------------------------------|
| Give access                 | GRANT                        | GRANT SELECT ON table<br>TO user;    |
| Take away<br>access         | REVOKE                       | REVOKE UPDATE ON table<br>FROM user; |
| Assign to<br>multiple users | Use comma-<br>separated list | TO user1, user2;                     |
| Manage by group             | Use ROLES                    | GRANT role_name TO<br>user;          |

# Transaction Control Language (TCL)

## 1. What is the purpose of the COMMIT and ROLLBACK commands in SQL?

In SQL, **COMMIT** and **ROLLBACK** are **transaction control commands** that manage changes made by a transaction in a database. Their primary purpose is to ensure data integrity and consistency.

---

### ◆ **COMMIT**

- **Purpose:** Saves all changes made during the current transaction to the database **permanently**.
- Once committed, the changes cannot be undone.
- Used when all operations in a transaction are successful and you want to make them permanent.

### **Example:**

```
BEGIN TRANSACTION;
```

```
UPDATE accounts SET balance = balance - 100 WHERE  
account_id = 1;
```

```
UPDATE accounts SET balance = balance + 100 WHERE  
account_id = 2;
```

```
COMMIT;
```

---

### ◆ **ROLLBACK**

- **Purpose:** Undoes all changes made during the current transaction.
- Used when something goes wrong during the transaction and you want to revert the database to its previous state.

**Example:**

```
BEGIN TRANSACTION;
```

```
UPDATE accounts SET balance = balance - 100 WHERE
account_id = 1;
```

-- Suppose an error occurs here

```
ROLLBACK;
```

---

## Summary

| Command Effect                                | Reversible? |
|-----------------------------------------------|-------------|
| COMMIT Makes changes permanent                | ✗ No        |
| ROLLBACK Reverts changes since the last BEGIN | ✓ Yes       |

These commands are essential in maintaining **atomicity** in ACID-compliant database transactions.

## 2. Explain how transactions are managed in SQL databases.

In SQL databases, **transactions** are used to group a sequence of operations into a single, logical unit of work that must be **entirely completed** or **entirely failed**. This ensures **data integrity** and follows the **ACID properties**.

---

## ◆ What Is a Transaction?

A **transaction** is a set of SQL operations (like INSERT, UPDATE, DELETE) that are executed as a single unit. A transaction **begins**, performs one or more operations, and then **either commits** or **rolls back** based on success or failure.

---

## ◆ ACID Properties of Transactions

1. **Atomicity** – All operations in a transaction are completed, or none are.
  2. **Consistency** – The database remains in a valid state before and after the transaction.
  3. **Isolation** – Concurrent transactions do not interfere with each other.
  4. **Durability** – Once committed, changes persist even after a system failure.
- 

## ◆ Transaction Lifecycle

1. **BEGIN** (or implicit start)

The transaction begins. Depending on the database, it can be:

- Implicit: Automatically starts with the first SQL statement.
- Explicit: Starts with BEGIN TRANSACTION.

## 2. SQL Operations

You perform one or more operations (e.g., INSERT, UPDATE, DELETE).

## 3. COMMIT

If all operations succeed, the transaction is committed, and changes are permanently saved.

## 4. ROLLBACK

If an error occurs, the transaction is rolled back, and all changes are undone.

---

### ◆ Example

```
BEGIN TRANSACTION;
```

```
UPDATE accounts SET balance = balance - 500 WHERE  
account_id = 1;
```

```
UPDATE accounts SET balance = balance + 500 WHERE  
account_id = 2;
```

```
-- If no errors, make it permanent
```

```
COMMIT;
```

```
-- If an error occurred:
```

```
-- ROLLBACK;
```

---

### ◆ Isolation Levels

SQL databases support different **isolation levels** to control how transactions interact:

- READ UNCOMMITTED
- READ COMMITTED
- REPEATABLE READ
- SERIALIZABLE

Each level balances **concurrency** and **data integrity** differently.

---

### ◆ Summary

Transactions in SQL databases ensure that operations are:

- Executed reliably
- Can be rolled back on failure
- Help maintain data consistency
- Are isolated from other users until committed

They are critical for applications involving **financial data**, **inventory systems**, or any environment where **data accuracy** is essential.

# SQL Joins

## 1. Explain the concept of JOIN in SQL. What is the difference between INNER JOIN, LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN?

In SQL, a **JOIN** is used to combine rows from two or more tables based on a **related column** between them — usually a foreign key. This is a powerful feature for retrieving data that spans multiple tables.

---

### ◆ Types of SQL JOINS

| Type                  | Description                                                                                               |
|-----------------------|-----------------------------------------------------------------------------------------------------------|
| INNER JOIN            | Returns only the matching rows from both tables.                                                          |
| LEFT JOIN             | Returns all rows from the <b>left</b> table, and matched rows from the right table (or NULL if no match). |
| RIGHT JOIN            | Returns all rows from the <b>right</b> table, and matched rows from the left table (or NULL if no match). |
| FULL<br>OUTER<br>JOIN | Returns all rows from both tables; unmatched rows get NULLs for missing values.                           |

---

### ◆ 1. INNER JOIN

Returns only rows that have matching values in **both** tables.

```
SELECT a.name, b.order_id
```

FROM customers a

INNER JOIN orders b ON a.customer\_id = b.customer\_id;

- ✓ Returns: Only customers who **have placed orders**
- 

#### ◆ 2. LEFT JOIN (or LEFT OUTER JOIN)

Returns all rows from the **left** table, plus matching rows from the right table. If no match exists, NULL is returned for right table columns.

SELECT a.name, b.order\_id

FROM customers a

LEFT JOIN orders b ON a.customer\_id = b.customer\_id;

- ✓ Returns: **All customers**, with their orders if any
  - ✗ No order → order\_id will be NULL
- 

#### ◆ 3. RIGHT JOIN (or RIGHT OUTER JOIN)

Returns all rows from the **right** table, plus matching rows from the left table. If no match exists, NULL is returned for left table columns.

SELECT a.name, b.order\_id

FROM customers a

RIGHT JOIN orders b ON a.customer\_id = b.customer\_id;

- ✓ Returns: **All orders**, with customer info if available
  - ✗ No customer → name will be NULL
-

## ◆ 4. FULL OUTER JOIN

Returns all rows when there is a match in **either** table.

Unmatched rows in either table will have NULL for the other side.

```
SELECT a.name, b.order_id
```

```
FROM customers a
```

```
FULL OUTER JOIN orders b ON a.customer_id = b.customer_id;
```

 Returns:

- Customers with orders
  - Customers **without** orders
  - Orders **without** customer data
- 

## ◆ Visual Summary

| Type            | Included Rows                                |
|-----------------|----------------------------------------------|
| INNER JOIN      | Matching rows only                           |
| LEFT JOIN       | All left + matching right (NULL if no match) |
| RIGHT JOIN      | All right + matching left (NULL if no match) |
| FULL OUTER JOIN | All left + all right (NULLs where no match)  |

---

## 2. How are joins used to combine data from multiple tables?

Joins are used in SQL to **combine data from multiple tables** based on a **related column**, usually a foreign key. This allows you to retrieve comprehensive results that span across several tables, like combining a customer's name with their order details or product information with supplier data.

---

### ◆ How Joins Work

Most relational databases are **normalized** — meaning data is split across multiple tables to reduce redundancy. A **JOIN** allows you to **reconstruct relationships** between these tables by combining rows where values in a specific column match.

---

### ◆ Step-by-Step Example

Imagine two tables:

#### **customers**

##### **customer\_id name**

1 Alice

2 Bob

#### **orders**

##### **order\_id customer\_id amount**

|     |   |     |
|-----|---|-----|
| 101 | 1 | 250 |
|-----|---|-----|

---

**order\_id customer\_id amount**

|     |   |     |
|-----|---|-----|
| 102 | 2 | 300 |
| 103 | 1 | 150 |

---

**◆ Example: INNER JOIN**

```
SELECT customers.name, orders.order_id, orders.amount
```

```
FROM customers
```

```
INNER JOIN orders ON customers.customer_id =  
orders.customer_id;
```

**Result:****name order\_id amount**

|       |     |     |
|-------|-----|-----|
| Alice | 101 | 250 |
| Bob   | 102 | 300 |
| Alice | 103 | 150 |

**Explanation:**

This joins the two tables where `customers.customer_id = orders.customer_id`, returning rows with **matching values**.

---

**◆ Other Join Types for Combining Data**

1. **LEFT JOIN**: Show all customers and any orders they have.
2. **RIGHT JOIN**: Show all orders and any associated customers.
3. **FULL OUTER JOIN**: Show all customers and all orders, matched where possible.

---

## ◆ Real-Life Use Cases

| Use Case                                                              | Tables Involved                | Join Used      |
|-----------------------------------------------------------------------|--------------------------------|----------------|
| Show employee names with their department names                       | employees, departments         | INNER JOIN     |
| Show all products and their supplier names (even if not supplied yet) | products, suppliers            | LEFT JOIN      |
| Show all sales and any customers involved                             | sales, customers               | RIGHT JOIN     |
| Combine students and courses to show all enrollments and open seats   | students, courses, enrollments | Multiple JOINs |

---

## Summary

- **Joins link tables** by matching columns (often foreign keys).
- Let you query **related data** as if it were in a single table.
- Essential for working with **normalized relational databases**.

# SQL Group By

## 1. What is the GROUP BY clause in SQL? How is it used with aggregate functions?

The GROUP BY clause in **SQL** is used to arrange **identical data into groups**. It's commonly used in combination with **aggregate functions** to perform operations (like SUM, COUNT, AVG, MAX, MIN) on each group of data.

---

### ◆ Syntax:

```
SELECT column_name, AGGREGATE_FUNCTION(column_name)  
FROM table_name  
GROUP BY column_name;
```

---

### ◆ How It Works:

- The GROUP BY clause groups rows that have the same values in specified columns into summary rows.
  - Then you can apply an aggregate function to each group.
- 

### ◆ Example:

Suppose you have a table called Sales:

**Region SalesAmount**

East 100

## **Region SalesAmount**

West 200

East 150

West 100

North 300

---

◆ **Query:**

```
SELECT Region, SUM(SalesAmount) AS TotalSales  
FROM Sales  
GROUP BY Region;
```

---

◆ **Result:**

## **Region TotalSales**

East 250

West 300

North 300

This query:

- Groups the rows by the Region column.
  - Calculates the SUM of SalesAmount for each group.
- 

◆ **Notes:**

- All columns in the SELECT list that are **not** aggregate functions **must** appear in the GROUP BY clause.
- You can group by **multiple columns**.

```
SELECT Region, Product, COUNT(*) AS SalesCount  
FROM Sales  
GROUP BY Region, Product;
```

---

## 2. Explain the difference between GROUP BY and ORDER BY.

The GROUP BY and ORDER BY clauses in SQL serve **different purposes**, though they are often used together. Here's a clear breakdown of their differences:

---

### ◆ 1. Purpose

#### Clause   Purpose

GROUP BY   Groups rows that have the same values in specified columns into **summary rows** (typically used with aggregate functions like SUM, AVG, COUNT).

ORDER BY   Sorts the result set based on one or more columns, in **ascending (ASC) or descending (DESC)** order.

---

### ◆ 2. Usage Example

Given a table Orders:

### **Customer Amount**

Alice      100

Bob      200

Alice      150

Bob      50

---

### **GROUP BY Example:**

```
SELECT Customer, SUM(Amount) AS TotalSpent
```

```
FROM Orders
```

```
GROUP BY Customer;
```

### **Result:**

### **Customer TotalSpent**

Alice      250

Bob      250

- Here, data is **grouped** by Customer, and aggregate function SUM() is applied.
- 

### **ORDER BY Example:**

```
SELECT Customer, Amount
```

```
FROM Orders
```

```
ORDER BY Amount DESC;
```

## Result:

### Customer Amount

Bob        200

Alice      150

Alice      100

Bob        50

- Here, rows are just **sorted** by Amount, not grouped or aggregated.
- 

### ◆ 3. Used With Aggregate Functions?

- GROUP BY → **Yes**, usually used with aggregate functions.
  - ORDER BY → **No**, but it can sort the output **after** aggregation.
- 

### ◆ 4. Order of Execution

In SQL query execution:

1. FROM
2. WHERE
3. GROUP BY
4. HAVING
5. SELECT
6. ORDER BY

This means GROUP BY happens **before** ORDER BY.

---

## ◆ 5. Can They Be Used Together?

 Yes! Example:

```
SELECT Customer, SUM(Amount) AS TotalSpent
```

```
FROM Orders
```

```
GROUP BY Customer
```

```
ORDER BY TotalSpent DESC;
```

This groups the data **and** sorts the result by the total amount spent.

---

# SQL Stored Procedure

## 1. What is a stored procedure in SQL, and how does it differ from a standard SQL query?

A **stored procedure** in SQL is a **precompiled collection of one or more SQL statements** that are stored on the database server and can be executed as a single unit. It is used to **encapsulate logic, improve performance, and simplify repetitive tasks**.

---

### ◆ What Is a Stored Procedure?

- It's like a **function** in programming languages but for SQL.
  - Written once, stored in the database, and can be called repeatedly.
  - Can accept **parameters** (input, output, or both).
  - Can include **control-flow logic** (e.g., IF, WHILE, BEGIN...END).
  - Supports error handling.
- 

### ◆ Basic Syntax (SQL Server Example):

```
CREATE PROCEDURE GetCustomerOrders  
    @CustomerID INT  
AS  
BEGIN  
    SELECT * FROM Orders WHERE CustomerID = @CustomerID;
```

END;

To run it:

```
EXEC GetCustomerOrders @CustomerID = 101;
```

---

#### ◆ Stored Procedure vs Standard SQL Query

| Feature                   | Stored Procedure                         | Standard SQL Query                                |
|---------------------------|------------------------------------------|---------------------------------------------------|
| <b>Definition</b>         | Named, stored set of SQL statements      | Single query (e.g., SELECT, UPDATE)               |
| <b>Reusability</b>        | Reusable; can be called many times       | One-time execution unless manually reused         |
| <b>Parameters</b>         | Can take input/output parameters         | Usually does not support parameters               |
| <b>Logic/Flow Control</b> | Can include complex logic (IF, LOOP)     | Usually just a single statement                   |
| <b>Performance</b>        | Compiled and cached for faster execution | Parsed and executed at runtime                    |
| <b>Security</b>           | Access can be controlled via permissions | Less granular control                             |
| <b>Maintenance</b>        | Centralized and easier to update logic   | Repeated logic must be updated in multiple places |

---

#### ◆ Example Use Cases

- Validating user input.
  - Performing multiple SQL operations in a transaction.
  - Automating reports or daily processes.
  - Encapsulating business logic (e.g., updating stock after a sale).
- 

## **2. Explain the advantages of using stored procedures.**

Stored procedures offer several **key advantages** that make them highly useful in SQL development and database management.

Here's a breakdown of their main benefits:

---

### **1. Improved Performance**

- Stored procedures are **precompiled** and stored in the database.
  - When executed, the database reuses the compiled execution plan, making them faster than sending raw SQL each time.
- 

### **2. Reusability and Maintainability**

- Write once, use many times.

- Centralized logic: if business rules change, update the procedure in one place rather than modifying multiple queries across applications.
- 

### 3. Security and Access Control

- Users can be granted permission to **execute** the procedure without giving direct access to the underlying tables.
  - Helps protect sensitive data and control operations more tightly.
- 

### 4. Reduced Network Traffic

- Instead of sending multiple SQL commands over the network, a single call to a stored procedure can perform many operations.
  - This reduces bandwidth usage and round trips between application and database.
- 

### 5. Modularity

- Logic can be broken into **logical blocks** using procedures, making complex operations easier to understand, maintain, and debug.
- 

### 6. Supports Complex Logic

- Stored procedures can include:
  - Conditional logic (IF, CASE)

- Loops (WHILE, FOR)
  - Error handling (TRY...CATCH or EXCEPTION)
  - Makes it easier to handle business rules that require multiple steps or checks.
- 

## 7. Improved Consistency

- Ensures that business logic is implemented the **same way every time**, avoiding duplication and inconsistencies.
- 

## 8. Transactional Control

- You can include BEGIN TRANSACTION, COMMIT, and ROLLBACK to ensure data integrity within a stored procedure.
- 

## 9. Abstract Complexity from Clients

- Application developers can call a stored procedure without needing to know the detailed SQL logic or table structure.
- 

## Summary Table:

| Benefit     | Description                 |
|-------------|-----------------------------|
| Performance | Precompiled and cached      |
| Reusability | Write once, use many times  |
| Security    | Fine-grained access control |

| <b>Benefit</b>  | <b>Description</b>                          |
|-----------------|---------------------------------------------|
| Efficiency      | Less network traffic                        |
| Maintainability | Easier updates to business logic            |
| Logic Support   | Can include flow control and error handling |
| Consistency     | Uniform business rule enforcement           |

---

# SQL View

## 1. What is a view in SQL, and how is it different from a table?

A **view** in SQL is a **virtual table** that is based on the result of a SQL SELECT query. Unlike a regular table, a view **does not store data itself**; it presents data from one or more tables dynamically whenever the view is queried.

---

### ◆ What Is a View?

- A **view** is a **named query** stored in the database.
  - It behaves like a table when queried, but it pulls data from underlying tables in real time.
  - Think of it as a **window** into your data.
- 

### ◆ Syntax Example:

```
CREATE VIEW ActiveCustomers AS  
SELECT CustomerID, Name, Email  
FROM Customers  
WHERE Status = 'Active';
```

Now, you can query it like a table:

```
SELECT * FROM ActiveCustomers;
```

---

### ◆ View vs Table: Key Differences

| <b>Feature</b>     | <b>View</b>                                            | <b>Table</b>                      |
|--------------------|--------------------------------------------------------|-----------------------------------|
| <b>Definition</b>  | Virtual table based on a SELECT query                  | Physical storage of data          |
| <b>Storage</b>     | Does <b>not</b> store data (except materialized views) | <b>Stores</b> data                |
| <b>Data Source</b> | Pulls data from one or more tables                     | Is the data source itself         |
| <b>Updatable?</b>  | Sometimes (read-only if complex joins used)            | Always                            |
| <b>Use Case</b>    | Simplify complex queries, security, abstraction        | Store raw data                    |
| <b>Maintenance</b> | Reflects live data from underlying tables              | Must be updated via INSERT/UPDATE |

---

#### ◆ Advantages of Views

- **Simplicity:** Hide complex joins and logic behind a simple name.
  - **Security:** Limit access to specific rows or columns of a table.
  - **Reusability:** Define once, use many times.
  - **Consistency:** Ensures consistent business logic.
- 

#### ◆ When to Use Views

- You want to **abstract** complex SQL logic.
  - You want to **control access** to certain data.
  - You need a **custom presentation** of the data.
- 

## 2. Explain the advantages of using views in SQL databases.

Using **views** in SQL databases offers several key advantages, especially when it comes to simplifying, securing, and managing data access. Here's a detailed look at the benefits:

---

### 1. Data Abstraction & Simplification

- Views allow you to **encapsulate complex SQL queries** behind a simple name.
  - Users can interact with a simple SELECT from a view instead of writing joins, filters, or calculations every time.
- ◆ *Example:* A view can hide a JOIN between multiple tables, presenting just the needed columns.
- 

### 2. Improved Security

- Views can **limit access** to sensitive data.
- You can expose only specific columns or rows to certain users without giving them direct access to the underlying tables.

- ◆ *Example:* A View can show only active customers or exclude salary fields from an employee table.
- 

### 3. Code Reusability & Maintainability

- Views promote **DRY (Don't Repeat Yourself)** by letting you reuse queries without rewriting them.
  - If the underlying logic changes, you can **update the view** once instead of updating every client query.
- 

### 4. Logical Data Independence

- Views provide a **layer of abstraction** from the underlying database structure.
  - You can change table structures without affecting users or applications if they query through views.
- 

### 5. Centralized Business Logic

- Business rules (like filters, calculations, or joins) can be embedded into views.
  - This ensures **consistency** and reduces the chance of logic duplication or errors in application code.
- 

### 6. Improved Query Readability

- Views can give **meaningful names** to calculated fields or complicated expressions, improving readability and clarity.

---

## 7. Support for Data Aggregation

- You can create views that include aggregates (e.g., totals, averages) for use in dashboards or reporting tools.
- 

## 8. Modular Query Design

- Break down large, complex queries into smaller pieces using views.
  - You can even build views on top of other views for layered logic.
- 

### Summary Table:

| Advantage            | Description                                    |
|----------------------|------------------------------------------------|
| Simplification       | Hides complex queries and joins                |
| Security             | Restricts access to sensitive data             |
| Reusability          | Centralizes frequently-used queries            |
| Maintainability      | Changes to logic only need to be made once     |
| Logical Independence | Shields users from changes in table structures |
| Consistency          | Uniform business logic across applications     |

---

# SQL Triggers

## 1. What is a trigger in SQL? Describe its types and when they are used.

A trigger in SQL is a **special kind of stored procedure** that **automatically executes** (or "fires") in response to specific events on a table or view, such as INSERT, UPDATE, or DELETE.

Triggers are used to **enforce rules**, **audit changes**, or **automate tasks** in the database.

---

### ◆ What Is a Trigger?

- A trigger is **bound to a table or view**.
  - It runs **automatically** when a specified database event occurs.
  - Unlike procedures, **you don't call a trigger manually**—the DBMS handles that.
- 

### ◆ Syntax Example (SQL Server):

```
CREATE TRIGGER trg_AuditLog
ON Employees
AFTER INSERT
AS
BEGIN
    INSERT INTO AuditLog (Action, Timestamp)
```

```
VALUES ('New employee added', GETDATE());  
END;
```

---

## ◆ Types of Triggers

### 1. AFTER Trigger (aka FOR Trigger)

- Executes **after** the triggering SQL statement.
  - Most commonly used.
  - **Use case:** Logging, cascading operations, enforcing complex business rules.
- ◆ **Example:** After inserting a new order, update inventory.
- 

### 2. BEFORE Trigger

- Executes **before** the triggering SQL statement.
  - Mostly supported in **MySQL, PostgreSQL**, not SQL Server.
  - **Use case:** Validating or modifying data before insert/update.
- ◆ **Example:** Ensure salary is not negative before inserting employee.
- 

### 3. INSTEAD OF Trigger

- Replaces the triggering action.
- Commonly used on **views**, where direct INSERT, UPDATE, or DELETE is not allowed.
- **Use case:** Custom behavior for changes on views.

- ◆ **Example:** Handle an UPDATE on a view by updating multiple base tables.
- 

## 4. Row-Level vs Statement-Level Triggers

- **Row-level:** Executes once **per affected row**.
  - **Statement-level:** Executes once **per SQL statement**, regardless of number of rows.
  - ◆ Most RDBMSs allow both; PostgreSQL uses FOR EACH ROW and FOR EACH STATEMENT.
- 

### ◆ When to Use Triggers

| Use Case                        | Description                                             |
|---------------------------------|---------------------------------------------------------|
| <b>Auditing</b>                 | Track changes to data for compliance (e.g., AuditLog)   |
| <b>Validation</b>               | Enforce business rules before data modification         |
| <b>Cascading actions</b>        | Automatically update/dependant tables (e.g., inventory) |
| <b>Restricting operations</b>   | Prevent certain changes under conditions                |
| <b>Maintaining derived data</b> | Automatically update summary or total fields            |

---

### ◆ Example Use Case

## **Prevent deleting high-value orders:**

```
CREATE TRIGGER trg_NoDeleteHighOrders  
BEFORE DELETE ON Orders  
FOR EACH ROW  
WHEN (OLD.TotalAmount > 10000)  
BEGIN  
    SIGNAL SQLSTATE '45000' SET MESSAGE_TEXT = 'Cannot delete  
high-value orders.';  
END;
```

---

### **Summary**

#### **Feature      Description**

|            |                                                                        |
|------------|------------------------------------------------------------------------|
| Trigger    | Auto-executed procedure in response to DML events                      |
| Types      | AFTER, BEFORE, INSTEAD OF; row-level or statement-level                |
| Main Uses  | Auditing, validation, automation, data integrity                       |
| Limitation | Can make debugging difficult if overused or used without documentation |

---

## **2. Explain the difference between INSERT, UPDATE, and DELETE triggers.**

The **difference between INSERT, UPDATE, and DELETE triggers** lies in the **type of database operation** they respond to. All three are types of **DML (Data Manipulation Language)** triggers that automatically execute when specific actions occur on a table or view.

Here's a detailed comparison:

---

### **◆ 1. INSERT Trigger**

- **Fires when:** A new row is added to a table using INSERT.
- **Use cases:**
  - Logging insertions (audit trail).
  - Automatically populating additional fields.
  - Enforcing constraints or default values.
- ◆ **Example:** Log every new customer added.

```
CREATE TRIGGER trg_AfterInsertCustomer
```

```
AFTER INSERT ON Customers
```

```
FOR EACH ROW
```

```
BEGIN
```

```
    INSERT INTO AuditLog (Action, Time)
```

```
    VALUES ('New customer inserted', NOW());
```

```
END;
```

---

## ◆ 2. UPDATE Trigger

- **Fires when:** A row is modified using UPDATE.
- **Use cases:**
  - Track changes to sensitive data.
  - Enforce business logic (e.g., salary cannot be decreased).
  - Compare old vs. new values for reporting or validation.
- ◆ **Example:** Log when employee salary changes.

```
CREATE TRIGGER trg_SalaryChange
```

```
AFTER UPDATE ON Employees
```

```
FOR EACH ROW
```

```
WHEN (OLD.Salary != NEW.Salary)
```

```
BEGIN
```

```
    INSERT INTO SalaryAudit (EmpID, OldSalary, NewSalary,  
    ChangedOn)
```

```
    VALUES (OLD.ID, OLD.Salary, NEW.Salary, NOW());
```

```
END;
```

---

## ◆ 3. DELETE Trigger

- **Fires when:** A row is removed using DELETE.
- **Use cases:**
  - Prevent deletion under certain conditions.

- Archive deleted records.
- Audit deleted data (since once deleted, it's gone unless stored).

◆ **Example:** Archive orders before deletion.

```
CREATE TRIGGER trg_BeforeDeleteOrder
BEFORE DELETE ON Orders
FOR EACH ROW
BEGIN
    INSERT INTO OrderArchive (OrderID, OrderDate, TotalAmount)
    VALUES (OLD.OrderID, OLD.OrderDate, OLD.TotalAmount);
END;
```

---

◆ **Summary Table**

| <b>Trigger Type</b> | <b>Fires On Operation</b> | <b>Typical Use Cases</b>                  | <b>Access to Rows</b> |
|---------------------|---------------------------|-------------------------------------------|-----------------------|
| INSERT              | When new data is added    | Log inserts, auto-fill fields, validation | NEW only              |
| UPDATE              | When data is modified     | Audit changes, validate updates           | OLD and NEW           |
| DELETE              | When data is removed      | Archive data, prevent deletion, audit     | OLD only              |

---

# Introduction to PL/SQL

## 1. What is PL/SQL, and how does it extend SQL's capabilities?

**PL/SQL (Procedural Language/Structured Query Language)** is Oracle's **procedural extension to SQL**. It combines **SQL's data manipulation power** with the features of **procedural programming languages**, such as control structures, loops, variables, and exception handling.

---

### ◆ What Is PL/SQL?

- Developed by **Oracle Corporation**.
  - Used to write **stored procedures, functions, triggers, packages, and anonymous blocks**.
  - Fully integrated with SQL, allowing seamless execution of queries and DML statements within procedural code.
- 

### ◆ How PL/SQL Extends SQL's Capabilities

| Feature                          | SQL                        | PL/SQL Enhancement                           |
|----------------------------------|----------------------------|----------------------------------------------|
| <b>Procedural Logic</b>          | No (only declarative)      | Supports IF, LOOP, WHILE, CASE               |
| <b>Variables &amp; Constants</b> | Not supported              | Full support for declaring and using         |
| <b>Exception Handling</b>        | Limited (just error codes) | Robust error handling using EXCEPTION blocks |

| <b>Feature</b>             | <b>SQL</b>             | <b>PL/SQL Enhancement</b>                |
|----------------------------|------------------------|------------------------------------------|
| <b>Modularity</b>          | No reusable code units | Supports procedures, functions, packages |
| <b>Loops and Branching</b> | Not supported          | Supports FOR, WHILE, EXIT, GOTO          |
| <b>Triggers</b>            | Not available directly | Easily written using PL/SQL              |

---

◆ **Example: PL/SQL vs SQL**

 **SQL (Limited)**

```
UPDATE Employees SET Salary = Salary + 1000 WHERE
Department = 'Sales';
```

 **PL/SQL (With Logic)**

```
DECLARE
    v_bonus NUMBER := 1000;
BEGIN
    FOR emp IN (SELECT * FROM Employees WHERE Department =
'Sales') LOOP
        IF emp.Salary < 5000 THEN
            UPDATE Employees SET Salary = Salary + v_bonus WHERE
ID = emp.ID;
        END IF;
    END LOOP;
END;
```

---

## ◆ Key Features of PL/SQL

- **Block Structure:** Organized into DECLARE, BEGIN, EXCEPTION, END.
  - **Tight SQL Integration:** Easily embed SQL statements in procedural logic.
  - **Portability:** Runs on any Oracle-supported platform.
  - **High Performance:** Executes in the database engine; reduces network traffic.
  - **Security:** Stored code can hide business logic from users.
- 

## ◆ Common Uses

- Writing **stored procedures and functions**
  - Defining **triggers**
  - Creating **packages** (collections of procedures/functions)
  - Automating complex **business logic**
  - Implementing **data validation rules**
- 

## Summary

| Aspect   | SQL Only          | PL/SQL                                 |
|----------|-------------------|----------------------------------------|
| Style    | Declarative       | Procedural + Declarative               |
| Use Case | Single statements | Complex logic, looping, error handling |

| Aspect      | SQL Only | PL/SQL                                |
|-------------|----------|---------------------------------------|
| Performance | Good     | Better in batch processing            |
| Modularity  | No       | Yes (functions, procedures, packages) |

---

## 2. List and explain the benefits of using PL/SQL.

Here are the key **benefits of using PL/SQL** (Procedural Language/SQL), especially in Oracle databases:

---

### 1. Tight Integration with SQL

PL/SQL is built specifically to work with SQL. You can write SQL queries directly inside PL/SQL blocks, allowing **seamless database interaction**.

 Example: Easily query, insert, update, or delete data within a PL/SQL program.

---

### 2. Supports Procedural Constructs

Unlike standard SQL, PL/SQL supports **procedural features** such as:

- Variables and constants
- Control statements (IF, CASE)
- Loops (FOR, WHILE, LOOP)

- Modular code (PROCEDURES, FUNCTIONS)

 This enables you to implement complex business logic directly within the database.

---

### 3. Better Performance

PL/SQL reduces the number of **network round-trips** between the application and the database by bundling multiple SQL statements into a single block.

 This leads to **faster execution**, especially in bulk processing or loops.

---

### 4. Modularity and Reusability

You can define **procedures, functions, and packages** in PL/SQL, which can be reused throughout your applications.

 Makes code easier to maintain, debug, and organize.

---

### 5. Robust Error Handling

PL/SQL offers **structured exception handling** using BEGIN...EXCEPTION...END blocks.

 You can catch and respond to database and application errors gracefully.

---

### 6. Improved Security

Stored PL/SQL code can **restrict access** to underlying tables. Users can be given access to **procedures** instead of direct table access.

- 🔒 This helps protect sensitive data and control permissions more precisely.
- 

## ✓ 7. Portability

PL/SQL code is **portable** across all platforms supported by Oracle.

- /blue/carry You can move applications and logic from one system to another with minimal changes.
- 

## ✓ 8. Efficient Transaction Control

PL/SQL supports full **transaction management**, including COMMIT, ROLLBACK, and SAVEPOINT.

- /blue/refresh Ensures data integrity and safe rollback of failed operations.
- 

## ✓ 9. Scalability for Large Applications

With features like **packages**, **triggers**, **cursors**, and **collections**, PL/SQL scales well for **enterprise-level applications** with thousands of lines of business logic.

---

## ✓ 10. Built-in Functions and Packages

Oracle provides a large library of **built-in packages** (like DBMS\_OUTPUT, UTL\_FILE, DBMS\_SQL, etc.) to help with tasks such as:

- Output logging
  - File operations
  - Email handling
  - Dynamic SQL
- 

## **Summary Table**

| <b>Benefit</b>   | <b>Description</b>                                    |
|------------------|-------------------------------------------------------|
| SQL Integration  | Combines procedural logic with SQL seamlessly         |
| Procedural Logic | Enables conditions, loops, and variables              |
| Performance      | Reduces network traffic; improves batch execution     |
| Modularity       | Supports reusable functions, procedures, and packages |
| Error Handling   | Robust EXCEPTION block for managing runtime errors    |
| Security         | Fine-grained access control through stored programs   |
| Portability      | Runs across Oracle environments without changes       |

| <b>Benefit</b>      | <b>Description</b>                             |
|---------------------|------------------------------------------------|
| Transaction Control | Full support for COMMIT, ROLLBACK, SAVEPOINT   |
| Scalability         | Suitable for large-scale, complex applications |
| Built-in Utilities  | Rich standard packages for common tasks        |

---

# PL/SQL Control Structures

## 1. What are control structures in PL/SQL? Explain the IF-THEN and LOOP control structures.

In **PL/SQL**, **control structures** are used to **control the flow** of execution within a program. They allow you to make decisions (IF statements), repeat actions (LOOP statements), and manage execution logic just like in traditional programming languages.

---

### ◆ Types of Control Structures in PL/SQL:

1. **Conditional Control** – IF, IF-THEN-ELSE, CASE
  2. **Iterative Control (Loops)** – LOOP, WHILE, FOR
  3. **Sequential Control** – GOTO, EXIT, NULL
- 

### 1. IF-THEN Control Structure

Used to execute a block of code **only if a condition is true**.

#### ◆ Syntax:

IF condition THEN

  -- Statements to execute if condition is true

END IF;

#### ◆ Extended Forms:

- **IF-THEN-ELSE**

IF condition THEN

  -- If true

ELSE

-- If false

END IF;

- **IF-THEN-ELSIF**

IF condition1 THEN

-- Block 1

ELSIF condition2 THEN

-- Block 2

ELSE

-- Default block

END IF;

◆ **Example:**

DECLARE

v\_salary NUMBER := 4000;

BEGIN

IF v\_salary < 5000 THEN

DBMS\_OUTPUT.PUT\_LINE('Below average salary');

ELSE

DBMS\_OUTPUT.PUT\_LINE('Good salary');

END IF;

END;

---

## 2. LOOP Control Structure

Used to **repeat a block of code** multiple times.

### ◆ Types of Loops:

1. **Basic LOOP**

2. **WHILE LOOP**

3. **FOR LOOP**

---

### ◆ a) Basic LOOP

Executes the block **indefinitely** until an EXIT statement is encountered.

LOOP

-- Statements

EXIT WHEN condition;

END LOOP;

### ◆ Example:

DECLARE

i NUMBER := 1;

BEGIN

LOOP

DBMS\_OUTPUT.PUT\_LINE('Counter: ' || i);

i := i + 1;

EXIT WHEN i > 5;

END LOOP;

END;

---

◆ **b) WHILE LOOP**

Repeats as long as the condition is TRUE.

WHILE condition LOOP

-- Statements

END LOOP;

◆ **Example:**

DECLARE

i NUMBER := 1;

BEGIN

WHILE i <= 5 LOOP

DBMS\_OUTPUT.PUT\_LINE('i = ' || i);

i := i + 1;

END LOOP;

END;

---

◆ **c) FOR LOOP**

Executes a block a **specific number of times**.

FOR i IN [REVERSE] start..end LOOP

-- Statements

END LOOP;

◆ **Example:**

```
BEGIN  
    FOR i IN 1..5 LOOP  
        DBMS_OUTPUT.PUT_LINE('i = ' || i);  
    END LOOP;  
END;
```

---

◆ **Summary Table**

| <b>Control Structure</b> | <b>Description</b>                         | <b>When to Use</b>    |
|--------------------------|--------------------------------------------|-----------------------|
| IF-THEN                  | Executes code based on a condition         | Decision making       |
| LOOP                     | Repeats block until manually exited        | Unknown iterations    |
| WHILE LOOP               | Repeats block as long as condition is true | Pre-condition loops   |
| FOR LOOP                 | Repeats a fixed number of times            | Known iteration range |

---

## 2. How do control structures in PL/SQL help in writing complex queries?

Control structures in **PL/SQL** play a crucial role in enabling you to write **complex, intelligent, and dynamic SQL logic** that goes far beyond what standard SQL can do on its own. Here's how they help:

---

### How Control Structures Help in Writing Complex Queries

---

#### ◆ 1. Enable Decision-Making with Conditions

PL/SQL's IF-THEN, IF-THEN-ELSE, and CASE constructs allow you to **conditionally execute SQL statements** based on business rules.

 **Use Case:** Apply different discounts based on customer type or order amount.

```
IF order_amount > 1000 THEN  
    discount := 0.15;  
  
ELSE  
    discount := 0.05;  
  
END IF;
```

---

#### ◆ 2. Allow Repetition for Batch Processing

Looping structures (LOOP, FOR, WHILE) let you process multiple rows or perform repeated actions, like running calculations for each row in a cursor.

 **Use Case:** Update bonuses for all employees based on performance.

```
FOR emp IN (SELECT * FROM employees) LOOP
```

```
    UPDATE employees
```

```
        SET bonus = emp.salary * 0.1
```

```
        WHERE id = emp.id;
```

```
END LOOP;
```

---

### ◆ 3. Support Dynamic Query Execution

Control structures allow you to **construct and execute queries at runtime** using variables and conditions.

 **Use Case:** Run different queries based on user input or parameters.

```
IF user_type = 'admin' THEN
```

```
    OPEN cur FOR SELECT * FROM full_report;
```

```
ELSE
```

```
    OPEN cur FOR SELECT * FROM summary_report;
```

```
END IF;
```

---

### ◆ 4. Handle Exceptions Gracefully

With EXCEPTION blocks, you can **detect and handle errors** in SQL operations without crashing the program.

 **Use Case:** Retry a query or log an error when a transaction fails.

```
BEGIN
```

```
    DELETE FROM orders WHERE order_id = 123;
```

```
EXCEPTION
```

```
    WHEN NO_DATA_FOUND THEN
```

```
        DBMS_OUTPUT.PUT_LINE('Order not found.');
```

```
END;
```

---

#### ◆ **5. Automate Multi-Step Logic**

PL/SQL lets you **bundle multiple operations**—queries, conditions, and updates—into one coherent program.

 **Use Case:** Calculate tax, update inventory, and insert invoice—all in one block.

```
BEGIN
```

```
    -- calculate tax
```

```
    -- update inventory
```

```
    -- insert invoice
```

```
END;
```

---



#### **Summary: What Control Structures Let You Do**

| Capability            | Benefit for Complex Queries                     |
|-----------------------|-------------------------------------------------|
| Conditional Execution | Handle different logic paths based on data      |
| Iteration             | Process rows or tasks repeatedly                |
| Modular Logic         | Break logic into readable steps                 |
| Dynamic SQL           | Build and run custom queries based on variables |
| Error Handling        | Manage failures without stopping the process    |
| Transaction Control   | Use COMMIT and ROLLBACK conditionally           |



### Bottom Line:

Without control structures, SQL can only **describe what data to get or modify**—but with **PL/SQL**, you can describe **how, when, and why** to perform database operations.

# SQL Cursors

## 1. What is a cursor in PL/SQL? Explain the difference between implicit and explicit cursors.

In **PL/SQL**, a **cursor** is a pointer to the **context area** that stores the result set of a SQL query. It allows you to **fetch rows one at a time**, making it essential for handling multi-row queries in procedural code.

---

### ◆ What Is a Cursor in PL/SQL?

- When a SQL query (like SELECT) returns **multiple rows**, a cursor is used to **iterate over the result set**.
  - Cursors help you **process each row individually**, especially in loops.
- 

### ◆ Types of Cursors in PL/SQL

| Type                   | Description                                                                       |
|------------------------|-----------------------------------------------------------------------------------|
| <b>Implicit Cursor</b> | Automatically created by PL/SQL for single-row SELECT, INSERT, UPDATE, or DELETE. |
| <b>Explicit Cursor</b> | Manually declared and controlled by the programmer for multi-row queries.         |

---

### Implicit Cursor

- Created automatically by Oracle **whenever a DML statement or SELECT INTO** is executed.

- Accessible via the SQL cursor.

◆ **Example:**

```
BEGIN
    UPDATE Employees SET Salary = Salary + 500 WHERE
    Department = 'Sales';

    DBMS_OUTPUT.PUT_LINE(SQL%ROWCOUNT || ' row(s)
updated.);

END;
```

◆ **Cursor Attributes:**

- SQL%ROWCOUNT – number of rows affected
- SQL%FOUND – TRUE if at least one row affected
- SQL%NOTFOUND – opposite of FOUND
- SQL%ISOPEN – always FALSE for implicit cursors

 **Explicit Cursor**

- Used when a SELECT statement returns **more than one row** and you want to process them individually.
- Requires four steps:
  1. **Declare**
  2. **Open**
  3. **Fetch**
  4. **Close**

◆ **Example:**

```

DECLARE

    CURSOR emp_cursor IS SELECT FirstName, Salary FROM
Employees WHERE Department = 'Sales';

    v_name Employees.FirstName%TYPE;
    v_salary Employees.Salary%TYPE;

BEGIN

    OPEN emp_cursor;

    LOOP

        FETCH emp_cursor INTO v_name, v_salary;
        EXIT WHEN emp_cursor%NOTFOUND;
        DBMS_OUTPUT.PUT_LINE(v_name || ':' || v_salary);

    END LOOP;

    CLOSE emp_cursor;

END;

```

◆ **Cursor Attributes:**

- cursor\_name%ROWCOUNT
- cursor\_name%FOUND
- cursor\_name%NOTFOUND
- cursor\_name%ISOPEN

◆ **Key Differences: Implicit vs Explicit Cursors**

| <b>Feature</b>       | <b>Implicit Cursor</b>                 | <b>Explicit Cursor</b>                              |
|----------------------|----------------------------------------|-----------------------------------------------------|
| Created By           | Automatically by Oracle                | Manually by the developer                           |
| Used For             | Single-row queries, DML (INSERT, etc.) | Multi-row SELECT queries                            |
| Control              | Automatic                              | Fully controlled (OPEN, FETCH, CLOSE)               |
| Naming               | Always referred to as SQL              | Custom name (e.g., emp_cursor)                      |
| Attributes Supported | SQL%FOUND,<br>SQL%ROWCOUNT,<br>etc.    | cursor_name%FOUND,<br>cursor_name%ROWCOUNT,<br>etc. |
| Flexibility          | Limited                                | Highly flexible; can use in loops and logic         |

---

## Summary

- Use **implicit cursors** for simple DML operations or SELECT INTO statements.
  - Use **explicit cursors** when dealing with **multi-row queries** and you need to process each row individually.
-

## 2. When would you use an explicit cursor over an implicit one?

You would use an **explicit cursor** over an **implicit cursor** in PL/SQL when you need **more control** over how you handle and process a **multi-row result set**. Implicit cursors work well for single-row operations or simple DML, but explicit cursors shine in more complex scenarios.

---

### When to Use an Explicit Cursor

---

#### ◆ 1. When the Query Returns Multiple Rows

- Implicit cursors can only handle **single-row SELECT INTO** queries.
- If you need to **fetch and process each row individually**, use an explicit cursor.

 **Example:** Loop through employees in a department to apply logic per person.

---

#### ◆ 2. When You Need Row-by-Row Processing

- Explicit cursors allow you to **manually fetch** rows one at a time.
  - Useful for logic that depends on **data from each row**, such as conditionals or calculations.
-

### ◆ 3. When You Need Cursor Attributes for Specific Control

- Explicit cursors give you detailed runtime feedback via:
    - %FOUND
    - %NOTFOUND
    - %ROWCOUNT
    - %ISOPEN
  - This helps control loop exits or track progress precisely.
- 

### ◆ 4. When You Want to Use a Cursor in a Loop

- Explicit cursors can be used in **manual LOOP/FETCH structures** or **cursor FOR loops**, making iteration clean and manageable.

```
FOR emp IN emp_cursor LOOP
```

```
    -- Process each employee
```

```
END LOOP;
```

---

### ◆ 5. When Using Parameterized Queries

- You can create **parameterized explicit cursors** that behave like functions and allow flexible input at runtime.

```
CURSOR emp_cursor (dept_id NUMBER) IS
```

```
SELECT * FROM employees WHERE department_id = dept_id;
```

---

### ◆ 6. When You Need Better Debugging or Maintainability

- With explicit cursors, logic is **modular and easy to debug**.
  - You can isolate query logic, making complex processes more readable.
- 

## When NOT to Use an Explicit Cursor

Avoid explicit cursors when:

- You're just doing a simple DML statement (INSERT, UPDATE, DELETE).
- You're working with a single-row SELECT INTO.

In those cases, **implicit cursors are simpler and more efficient**.

---

## Summary Table

| Scenario                           | Use Explicit Cursor?                                                                                      |
|------------------------------------|-----------------------------------------------------------------------------------------------------------|
| Fetching multiple rows             |  Yes                   |
| Need to process rows individually  |  Yes                   |
| Using SELECT INTO for one row      |  No                    |
| Just tracking row count from DML   |  No (use SQL%ROWCOUNT) |
| Need parameterized, reusable query |  Yes                   |

---

## Rollback and Commit Save point

**1. Explain the concept of SAVEPOINT in transaction management. How do ROLLBACK and COMMIT interact with save points?**

◆ **What is a SAVEPOINT in Transaction Management?**

A **SAVEPOINT** is a marker within a **transaction** that allows you to **partially roll back** to a specific point **without undoing the entire transaction**.

It acts like a **checkpoint** inside a transaction.

---

 **Key Benefits of SAVEPOINT**

- Allows **fine-grained control** within a transaction.
  - Lets you **undo part of the work** without affecting the rest.
  - Useful for complex transactions with **multiple logical units of work**.
- 

◆ **Syntax**

SAVEPOINT savepoint\_name;

ROLLBACK TO savepoint\_name;

---

◆ **Example:**

BEGIN;

-- Step 1

```
INSERT INTO Employees VALUES (101, 'Alice', 5000);
```

```
SAVEPOINT after_alice;
```

-- Step 2

```
INSERT INTO Employees VALUES (102, 'Bob', 6000);
```

```
SAVEPOINT after_bob;
```

-- Step 3

```
INSERT INTO Employees VALUES (103, 'Charlie', 7000);
```

-- Something goes wrong, rollback to after Alice

```
ROLLBACK TO after_alice;
```

-- Only Alice's insert remains

```
COMMIT;
```

---

◆ **How ROLLBACK Works with SAVEPOINT**

| Command                        | Effect                                                                   |
|--------------------------------|--------------------------------------------------------------------------|
| ROLLBACK;                      | Undoes <b>entire transaction</b> , including all changes and savepoints. |
| ROLLBACK TO<br>savepoint_name; | Undoes all changes <b>after the savepoint</b> , but keeps earlier ones.  |

---

#### ◆ How COMMIT Interacts with SAVEPOINT

- COMMIT **ends the transaction** and **makes all changes permanent**, including those before and after any SAVEPOINT.
  - After COMMIT, all savepoints are **cleared**—you cannot roll back to them anymore.
- 



#### Summary Table

| Feature        | Description                                         |
|----------------|-----------------------------------------------------|
| SAVEPOINT      | Creates a named point to partially roll back to     |
| ROLLBACK       | Undoes all changes since transaction start          |
| ROLLBACK<br>TO | Undoes changes only up to a specific SAVEPOINT      |
| COMMIT         | Makes all changes permanent; removes all savepoints |

---



#### Real-World Use Case

Imagine a banking system where you:

1. Debit one account
2. Credit another
3. Log the transaction

If the logging fails, you might want to ROLLBACK TO the point before logging, without reversing the debit/credit.

---

## 2. When is it useful to use save points in a database transaction?

Using **SAVEPOINT** in a database transaction is particularly useful when you need **fine-grained control over what gets rolled back** without discarding the entire transaction. This is valuable in **complex or multi-step operations** where partial failure is acceptable or expected.

---

### When Is It Useful to Use SAVEPOINTS?

---

#### ◆ 1. Handling Partial Failures Gracefully

You may want to **undo only part of a transaction** if a specific step fails, but keep the rest.

 **Example:** In a multi-step order process (inventory update, payment, logging), if logging fails, rollback only the log step:

SAVEPOINT before\_log;

-- attempt logging

ROLLBACK TO before\_log; -- but keep inventory and payment steps

---

## ◆ 2. Complex Business Transactions

In workflows like **banking, e-commerce, or HR systems**, many interrelated changes occur. SAVEPOINTS let you **checkpoint progress** and **backtrack** only if needed.

 Example: Approving a loan involves credit checks, document verification, and scoring. You can rollback to different SAVEPOINTS depending on which part fails.

---

## ◆ 3. Nested Operations in Loops

When performing operations inside a loop (e.g., batch updates or inserts), SAVEPOINTS let you **rollback only the failed iteration**, not the whole batch.

 Useful for **bulk processing** or **data migration** scripts.

---

## ◆ 4. Testing and Debugging

When developing or debugging, SAVEPOINTS allow you to **experiment with partial changes** and roll back to specific points to test different outcomes.

---

## ◆ 5. User-Driven Logic

If your app lets users choose to **cancel or undo part of a transaction**, SAVEPOINTS let you implement that logic without discarding everything.

 Example: A form with multiple tabs where user cancels only Tab 3's data, not the entire form.

---

#### ◆ **6. Error Recovery Without Total Rollback**

In PL/SQL or stored procedures, you can catch exceptions and roll back only part of a transaction.

BEGIN

SAVEPOINT step1;

-- Do something

SAVEPOINT step2;

-- Do another step

EXCEPTION

WHEN OTHERS THEN

ROLLBACK TO step1; -- Go back to safe point

END;

---

#### **Summary: When to Use SAVEPOINTS**

| Scenario                | Benefit                                    |
|-------------------------|--------------------------------------------|
| Multi-step transactions | Avoid losing all progress on small failure |

| Scenario                                 | Benefit                             |
|------------------------------------------|-------------------------------------|
| Loops with multiple inserts/updates      | Roll back just the failed iteration |
| Optional user actions                    | Allow undo of specific sections     |
| Complex business logic with dependencies | Control rollback granularity        |
| Error handling in procedures/triggers    | Provide fallback points             |

---

 **Bottom Line:**

Use SAVEPOINT when you need **control, flexibility, and fault tolerance** within transactions—especially in systems where **not everything must succeed to commit**.